APPONIX™
academy
SINCE 2013

Customer Churn Prediction and Analysis for Business Decision Making

Course: DATA ANALYTICS

Submitted By:

AYESHA MUSKAN A

TABLE OF CONTENTS

SL. NO	Contents	PAGE NO
01	Abstract	01
02	Introduction	02
03	Project Objectives	03
04	Tools and Techniques	04
05	Data Description	05
06	System Architecture	07
07	Project Methodology	08
08	Data Cleaning & Preprocessing	10
09	Exploratory Data Analysis 9.1 Churn Distribution 9.2 Monthly Charges vs churn 9.3 Tenure vs churn 9.4 Correlation Heatmap	10
10	Feature Engineering	12
11	Predictive Modeling	13
12	Model Evaluation 12.1 Confusion Matrix 12.2 ROC Curve	14
13	Dashboard Interpretation	15
14	Business Insights	17
15	Business Recommendations	18
16	Conclusion	19
17	Future Scope	20
18	References	20

Customer Churn Prediction and Analysis for Business Decision Making

1. ABSTRACT

Customer churn prediction is a critical problem in modern businesses, especially in industries such as telecommunications, banking, and subscription-based services. Churn refers to the phenomenon where customers stop using a company's product or service. Identifying customers who are likely to churn enables organizations to take proactive measures to retain them, thereby reducing revenue loss and improving customer satisfaction.

This project focuses on predicting customer churn using machine learning techniques. The dataset used contains customer-related information such as tenure, monthly charges, total charges, and other behavioral attributes. The data underwent preprocessing steps including handling missing values, converting data types, and feature engineering to improve model performance. Exploratory Data Analysis (EDA) was performed to understand patterns and relationships between different features and churn behavior.

Two classification models, Logistic Regression and Random Forest, were implemented to predict whether a customer will churn or not. The models were trained on preprocessed data and evaluated using performance metrics such as accuracy, confusion matrix, and Area Under the ROC Curve (AUC). The comparison of models helps in identifying the most effective approach for churn prediction.

The results demonstrate that machine learning models can effectively predict customer churn and highlight the key factors influencing customer retention. This project provides valuable insights that can help businesses make data-driven decisions and design strategies to improve customer loyalty.

2. INTRODUCTION

In today's competitive business environment, customer retention has become as important as customer acquisition. Companies invest significant resources in attracting new customers, but losing existing customers can lead to substantial financial losses. Customer churn is a major challenge faced by organizations, particularly in industries where customers have multiple alternatives. Therefore, predicting and reducing churn has become a key focus area for data-driven businesses.

Customer churn occurs when a customer decides to discontinue using a company's product or service. Understanding the reasons behind churn is essential for developing effective retention strategies. With the advancement of data analytics and machine learning, it is now possible to analyze historical customer data and predict future churn behavior with considerable accuracy.

This project aims to build a machine learning model that can predict customer churn based on various features such as tenure, billing information, and usage patterns. By analyzing these factors, the model helps in identifying customers who are at high risk of leaving. This allows businesses to take targeted actions, such as offering discounts or improving services, to retain valuable customers.

The project involves several stages, including data collection, data preprocessing, exploratory data analysis, feature engineering, model building, and evaluation. Visualization techniques are used to gain insights into the data, while machine learning algorithms are applied to develop predictive models. The performance of these models is assessed using standard evaluation metrics.

Overall, this project demonstrates how data analysis and machine learning can be leveraged to solve real-world business problems. It highlights the importance of understanding customer behavior and using predictive analytics to enhance decision-making and improve business outcomes.

3. PROJECT OBJECTIVES

The main objective of this project is to develop a machine learning-based system to predict customer churn and analyze the factors influencing customer retention. The project focuses on leveraging data analysis and predictive modeling techniques to identify customers who are likely to discontinue a service.

The specific objectives of this project are as follows:

- To analyze customer data and understand patterns related to churn behavior using exploratory data analysis techniques.
- To preprocess the dataset by handling missing values, converting data types, and preparing the data for model building.
- To perform feature engineering by creating new meaningful features that can improve the predictive performance of the model.
- To build and implement machine learning models such as Logistic Regression and Random Forest for churn prediction.
- To evaluate the performance of the models using metrics like accuracy, confusion matrix, and ROC-AUC score.
- To compare different models and identify the most effective model for predicting customer churn.
- To identify key features that significantly influence customer churn using feature importance analysis.
- To provide insights that can help businesses take proactive measures to reduce customer churn and improve retention strategies.

4. TOOLS & TECHNIQUES

This project utilizes various tools and techniques from the fields of data analysis, machine learning, and data visualization to effectively predict customer churn. The combination of these tools enables efficient data processing, model building, and insightful analysis.

Tools Used:

- **Python**: Python is the primary programming language used for this project due to its simplicity and powerful libraries for data analysis and machine learning.
- **VS Code**: Used as the development environment for writing, executing, and testing the code.
- **Power BI**: Used for creating interactive dashboards and visualizing business insights derived from the dataset.
- **SQL (Structured Query Language)**: Used for basic data querying and understanding structured data (if applicable in the project workflow).

Libraries and Frameworks:

- **Pandas**: Used for data manipulation, cleaning, and handling structured datasets.
- **NumPy**: Used for numerical computations and array operations.
- **Matplotlib**: Used for basic data visualization and plotting graphs.
- **Seaborn**: Used for advanced and more visually appealing statistical visualizations.
- **Scikit-learn**: Used for implementing machine learning models, data preprocessing, and evaluation metrics.

Techniques Used:

- **Data Preprocessing**: Includes handling missing values, converting data types, and preparing the dataset for analysis.
- **Exploratory Data Analysis (EDA)**: Used to analyze and visualize patterns, trends, and relationships within the data using plots such as count plots, box plots, and heatmaps.

- **Feature Engineering**: Creation of new features such as *Average Charges* to improve model performance and capture additional insights.
- **Data Scaling**: Standardization of features using StandardScaler to ensure better performance of machine learning models.
- **Machine Learning Algorithms**
    - Logistic Regression
    - Random Forest Classifier
- **Model Evaluation Techniques**
  Evaluation of model performance using:
    - Accuracy Score
    - Confusion Matrix
    - Classification Report
    - ROC Curve and AUC Score

5. DATA DESCRIPTION

The dataset used in this project is **Customer Churn Dataset**, which contains information about customers, their usage behavior, and whether they have churned or not. The dataset is used to analyze customer retention patterns and build predictive models for churn classification.

The dataset consists of multiple features related to customer demographics, account information, and billing details. Each row in the dataset represents a unique customer, while each column represents a specific attribute of that customer.

The target variable in this dataset is **Churn**, which indicates whether a customer has discontinued the service (1) or not (0). The remaining features act as independent variables used to predict churn behavior.

Dataset Overview:

- Total Records: (mention your dataset rows, e.g., 7000+)
- Total Features: (number of columns after cleaning)
- Target Variable: **Churn**
- Data Type: Mixed (Numerical and Categorical)

Feature Description

Feature Name	Description
customerID	Unique identifier for each customer (removed during preprocessing)
Tenure	Number of months the customer has stayed with the company
MonthlyCharges	Amount charged to the customer monthly
TotalCharges	Total amount charged to the customer
AvgCharges	Average charges per month (engineered feature)
Churn	Target variable indicating whether the customer left (1) or stayed (0)

Data Characteristics

- The dataset contains both numerical and categorical features.
- Some columns required data type conversion (e.g., TotalCharges).
- Missing values were present and handled using median imputation.
- Feature engineering was applied to create additional meaningful attributes.

6. SYSTEM ARCHITECTURE

The system architecture of the Customer Churn Prediction project represents the overall workflow and the sequence of steps involved in transforming raw data into meaningful predictions. It illustrates how data flows through different stages, including preprocessing, analysis, model building, and evaluation. The architecture begins with the input data, which consists of customer-related information such as tenure, billing details, and service usage. This raw data is first processed through a data preprocessing stage, where irrelevant columns are removed, missing values are handled, and data types are converted into appropriate formats for analysis.

After preprocessing, the data is passed to the Exploratory Data Analysis (EDA) stage. In this phase, various visualization techniques such as count plots, box plots, and heatmaps are used to understand patterns, trends, and relationships between different features and the target variable (Churn). The next stage involves feature engineering and data transformation, where new features such as average charges are created to enhance the predictive capability of the models. The processed data is then split into training and testing datasets, followed by feature scaling to normalize the data.

In the model building stage, machine learning algorithms such as Logistic Regression and Random Forest are applied to the training data. These models learn patterns from historical data and are then used to predict churn on unseen test data. The performance of the models is evaluated using metrics such as accuracy, confusion matrix, and ROC-AUC score. Finally, the output of the system includes predicted churn values, performance metrics, and visualizations such as confusion matrix and ROC curves. These results are further used to derive business insights and recommendations for improving customer retention.

System Workflow:

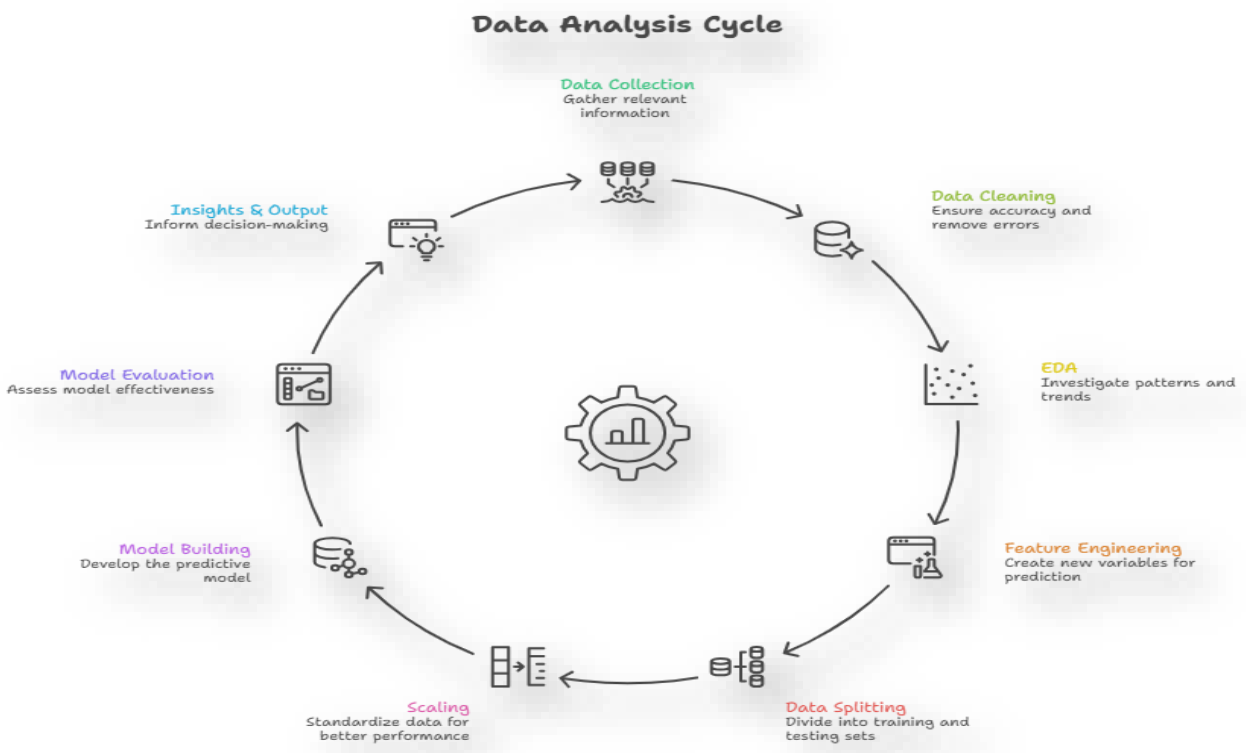

7. PROJECT METHODOLOGY

The project methodology outlines the systematic approach followed to develop the Customer Churn Prediction model. It includes all the steps from data collection to model evaluation, ensuring a structured and efficient workflow. The first step in the methodology is data collection, where the customer churn dataset is obtained and loaded into the system. The dataset contains various attributes related to customer behavior, billing information, and service usage, which are essential for analysis and prediction.

The next step involves data preprocessing, where the dataset is cleaned and prepared for analysis. This includes removing unnecessary columns such as customer identifiers, handling missing values, and converting data types into appropriate formats. These steps ensure that the data is consistent and suitable for further processing. Following preprocessing, Exploratory Data Analysis (EDA) is performed to understand the data distribution and relationships

between variables. Visualization techniques such as count plots, box plots, and heatmaps are used to identify patterns and trends related to customer churn. In the next stage, feature engineering is applied to create new meaningful features that can enhance model performance. For example, average charges are calculated to better represent customer spending behavior. The dataset is then split into training and testing sets to evaluate the model's performance on unseen data.

After data preparation, machine learning models such as Logistic Regression and Random Forest are implemented. These models are trained using the training dataset and then used to make predictions on the test dataset. Feature scaling is applied where necessary to improve model performance.

Finally, the models are evaluated using various performance metrics such as accuracy, confusion matrix, and ROC-AUC score. The results are analyzed to determine the effectiveness of each model, and insights are derived to support business decision-making.

Methodology Steps:

- **Data Collection**: Gathering the customer churn dataset for analysis and model building.
- **Data Preprocessing**: Cleaning the data by handling missing values, removing unnecessary columns, and converting data types.
- **Exploratory Data Analysis (EDA)**: Analyzing and visualizing data to identify patterns, trends, and relationships.
- **Feature Engineering**: Creating new meaningful features to improve model performance and insights.
- **Data Splitting & Scaling**: Dividing data into training and testing sets and standardizing features for better model accuracy.
- **Model Building**: Applying machine learning algorithms to train models for predicting customer churn.
- **Model Evaluation**: Assessing model performance using metrics like accuracy, confusion matrix, and ROC-AUC score.
- **Insights & Interpretation** – Deriving meaningful conclusions and business insights from the model results.

8. DATA CLEANING & PREPROCESSING

Data cleaning and preprocessing were performed to ensure the dataset is accurate, consistent, and suitable for machine learning models.

Steps Performed:

- Removed unnecessary column: customerID
- Converted TotalCharges from object to numeric format
- Handled missing values using median imputation
- Converted target variable Churn into integer format
- Reset dataset index for proper alignment
- Verified that no null values are present in the dataset

9. EXPLORATORY DATA ANALYSIS (EDA)

EDA was performed to understand the distribution of data and identify patterns related to customer churn using various visualizations.

Plot 1: Churn Distribution (Count Plot)

Description: This plot shows the distribution of customers who have churned and those who have not.

Insight: The dataset is slightly imbalanced, with a higher number of non-churn customers compared to churn customers.

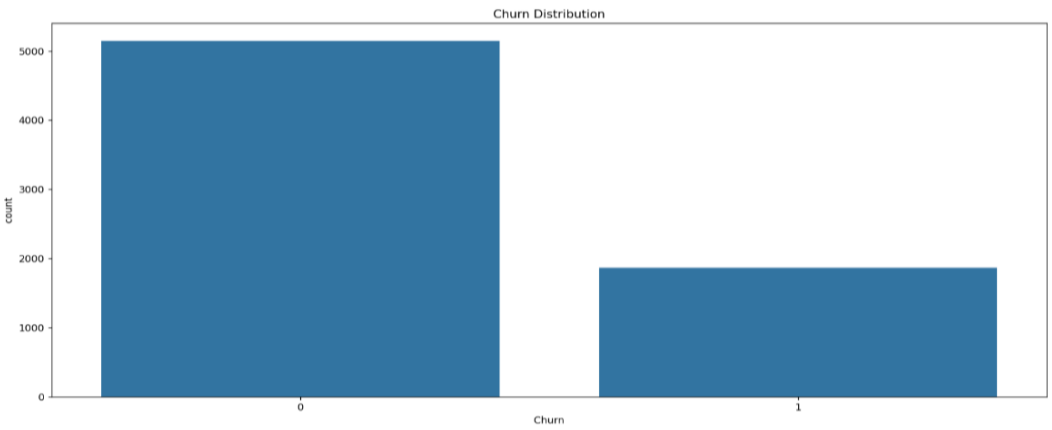

Plot 2: Monthly Charges vs Churn (Box Plot)

Description: This plot compares the distribution of monthly charges between churned and non-churned customers.

Insight: Customers with higher monthly charges tend to churn more frequently, indicating pricing may influence customer decisions.

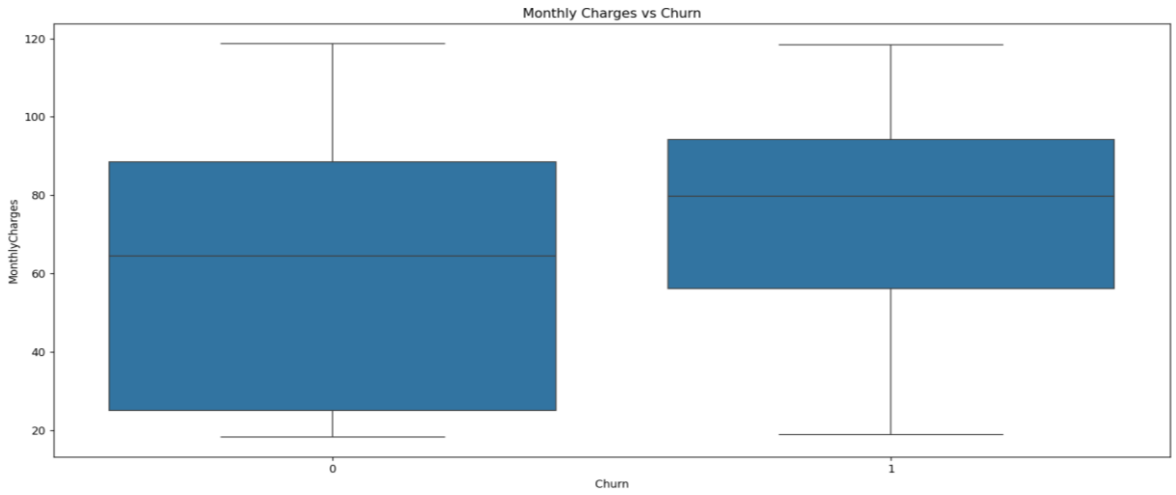

Plot 3: Tenure vs Churn (Box Plot)

Description: This plot shows the relationship between customer tenure and churn status.

Insight: Customers with lower tenure are more likely to churn, while long-term customers tend to stay with the service.

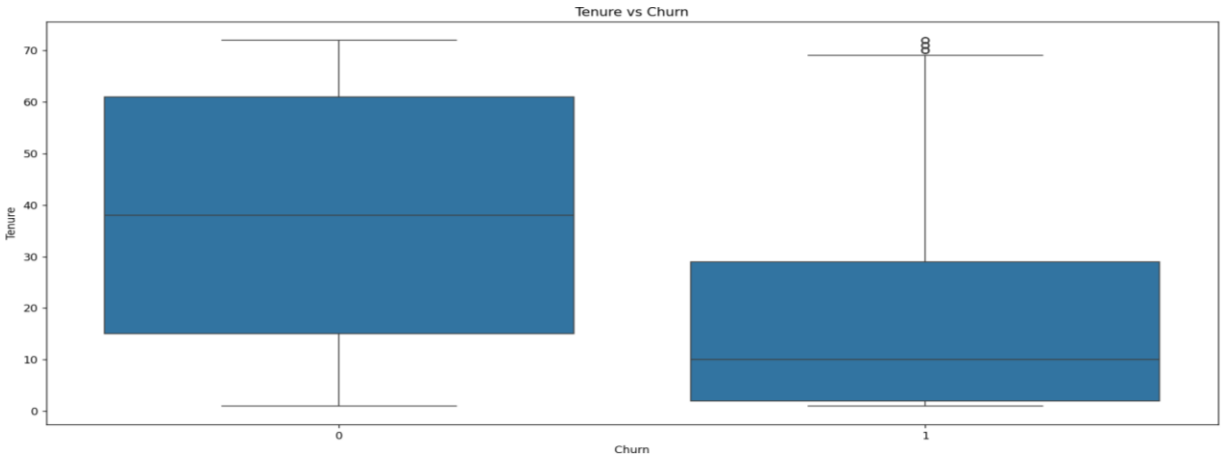

Plot 4: Correlation Heatmap

Description: The heatmap displays the correlation between numerical features in the dataset.

Insight: Features such as Monthly Charges and Total Charges show a strong relationship with churn, indicating their importance in prediction.

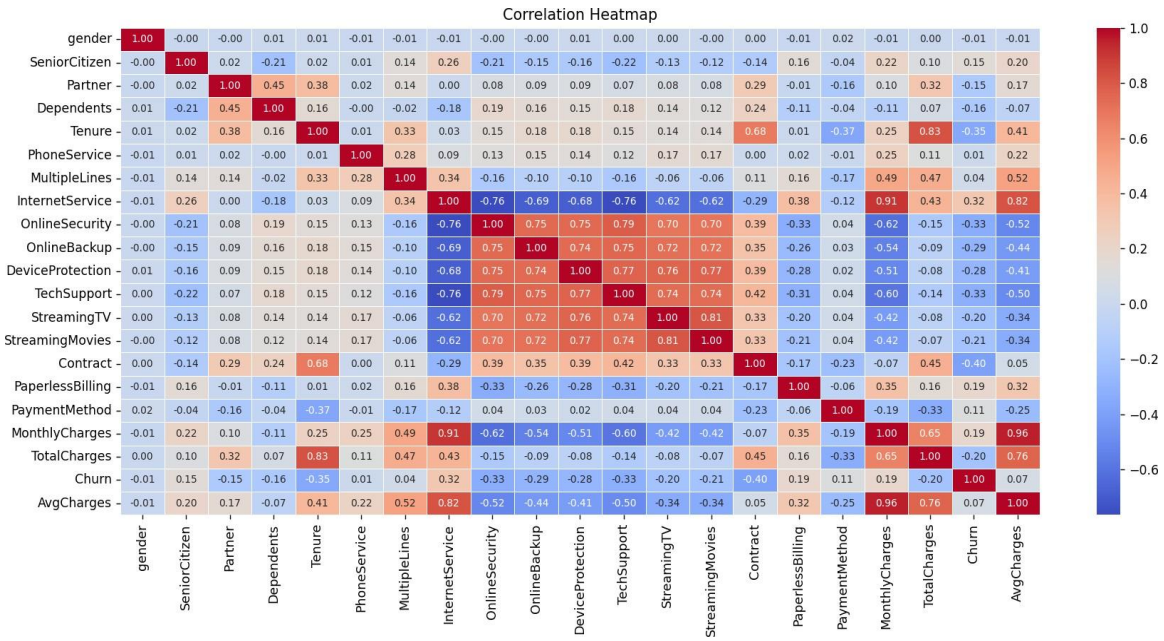

10. FEATURE ENGINEERING

Feature engineering is the process of creating new meaningful features from existing data to improve the performance of machine learning models. In this project, a new feature was derived to better represent customer spending behavior.

Feature Created:

- **AvgCharges (Average Charges)**
  This feature is calculated as the ratio of total charges to customer tenure:
- **AvgCharges = TotalCharges / (Tenure + 1)**

Purpose of Feature Engineering

- To capture the average spending pattern of customers.
- To provide a normalized view of customer charges over time.

- To enhance model performance by introducing a more informative feature.

Feature Importance (Bar Plot)

Description: This plot shows the importance of different features in predicting customer churn based on the Random Forest model.

Insight: Top features such as Monthly Charges, Tenure, and Total Charges have the highest influence on churn prediction.

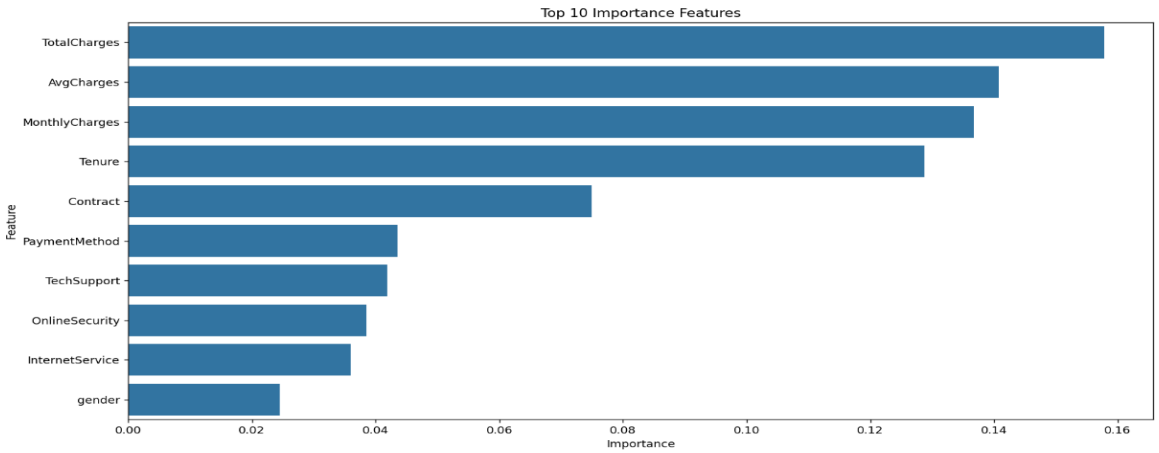

11. PREDICTIVE MODELING

Logistic Regression was used as a baseline classification model for predicting customer churn based on the relationship between input features and the target variable. Random Forest Classifier was used as an advanced ensemble model to capture complex patterns in the data and improve prediction accuracy.

The dataset was split into 80% training data and 20% testing data to train and evaluate the models effectively. Feature scaling was applied to standardize the input features before model training. Both models were used to predict customer churn, and their performance was compared to identify the most effective model for accurate prediction.

12. MODEL EVALUATION

The performance of the machine learning models was evaluated using different classification metrics to measure their effectiveness in predicting customer churn.

Evaluation Metrics Used:

- **Accuracy Score** : Measures the overall correctness of the model predictions.
- **Confusion Matrix**: Shows the number of correct and incorrect predictions made by the model.
- **Classification Report**: Provides precision, recall, and F1-score for better performance analysis.
- **ROC-AUC Score**: Evaluates the model's ability to distinguish between churn and non-churn customers.
- **ROC Curve**: Visual representation of model performance across different thresholds.

Model Comparison: Both Logistic Regression and Random Forest models were evaluated using the above metrics. Random Forest showed better performance in terms of accuracy and AUC score due to its ability to handle complex patterns in the data

Confusion Matrix (Heatmap)

Description: The confusion matrix visualizes the performance of the classification model by comparing actual and predicted values.

Insight: It shows how many customers were correctly and incorrectly classified, helping evaluate model accuracy and errors.

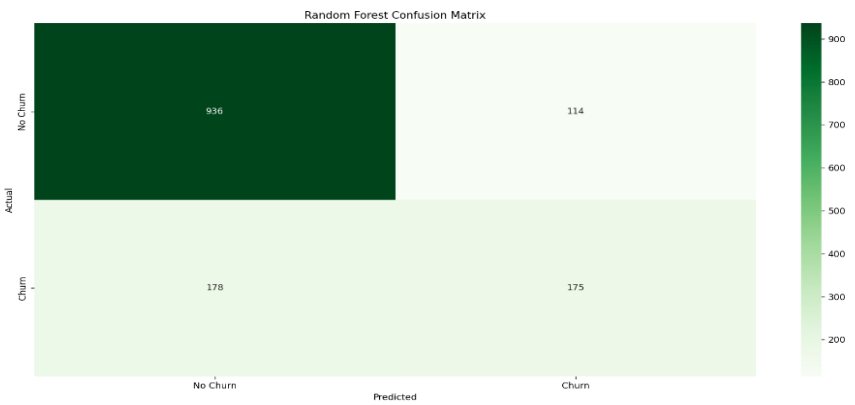

ROC Curve:

Description: The ROC curve plots the True Positive Rate against the False Positive Rate for different threshold values.

Insight: A curve closer to the top-left corner indicates better model performance, and higher AUC value represents a more accurate model

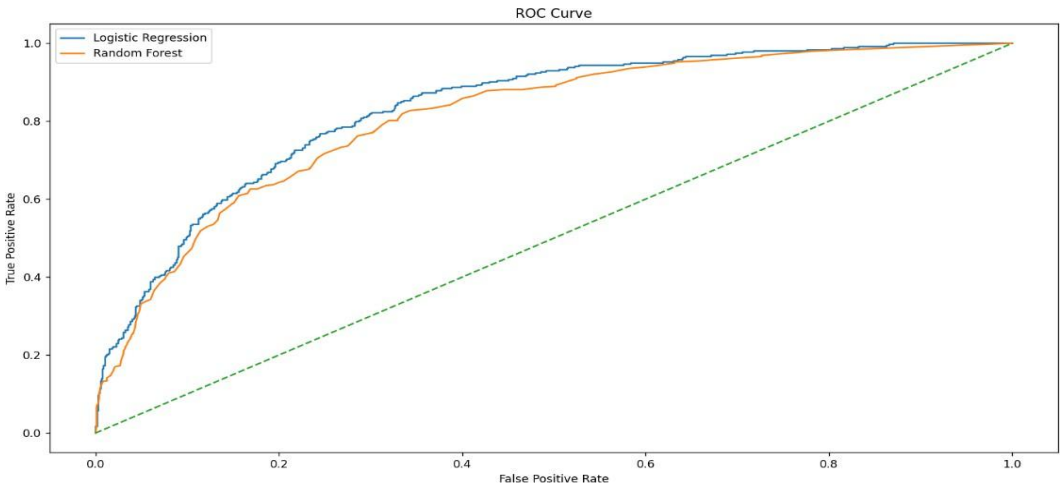

13. Dashboard Interpretation

This Customer Churn Analysis Dashboard provides a comprehensive overview of customer behavior, highlighting churn trends across contract types, payment methods, tenure groups, and charge categories. It helps identify key factors influencing customer retention and churn by visually comparing churned and active customers. The insights support data-driven decisions to improve customer retention strategies and reduce churn rate.

Key Performance Indicators (KPIs):

- Total Customers: **7011**
- Churn Customers: **1864**
- Active Customers: **3386**
- Average Monthly Charges: **64.80**
- Churn Rate: **26.59%**
- Average Tenure: **32.43 months**

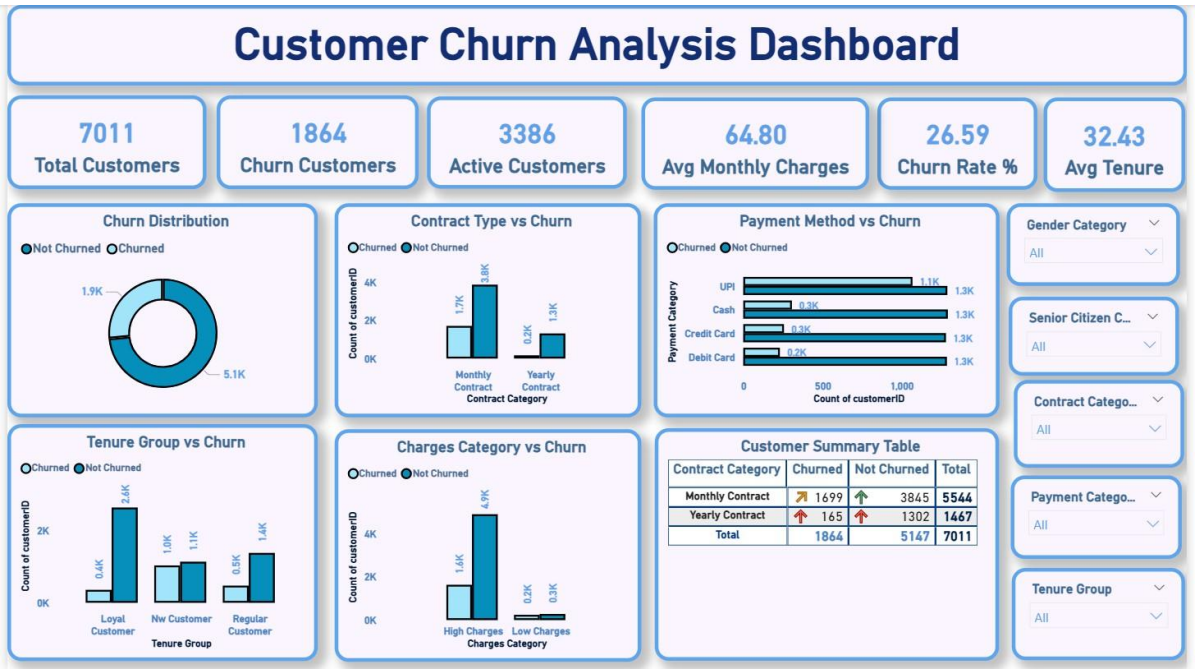

Cards (KPIs):

- **Total Customers:** Shows the total number of customers. It gives an overall view of the dataset size.
- **Churn Customers:** Shows how many customers left the service. It helps understand customer loss.
- **Active Customers:** Displays customers who are still using the service. It represents retained users.
- **Avg Monthly Charges:** Shows the average monthly bill of customers. It helps analyze pricing impact.
- **Churn Rate %:** Displays the percentage of customers who churned. It is a key performance metric.
- **Avg Tenure:** Shows the average time customers stay. It indicates customer loyalty.

Charts / Graphs:

- **Churn Distribution:** Shows churned vs non-churned customers. It gives a quick overall comparison.
- **Contract Type vs Churn:** Compares churn for monthly and yearly plans. It shows how contracts affect churn.

- **Payment Method vs Churn:** Shows churn based on payment types. It helps identify risky payment methods.
- **Tenure Group vs Churn:** Shows churn across different customer durations. It helps find which group churns more.
- **Charges Category vs Churn:** Compares churn for high and low charges. It shows pricing impact on churn.

Table:

- **Customer Summary Table:** Shows churn and non-churn data by contract type. It helps in detailed comparison.

Slcers (Filters):

- **Gender:** Filters data by gender. Helps compare churn between genders.
- **Senior Citizen:** Filters by senior citizen status. Helps analyze age-based churn.
- **Contract Type:** Filters by contract. Helps focus on specific plans.
- **Payment Method:** Filters by payment type. Helps analyze payment trends.
- **Tenure Group:** Filters by customer duration. Helps analyze retention over time.

14. BUSINESS INSIGHTS

The analysis of customer churn data provides several important business insights that can help organizations improve customer retention and reduce revenue loss.

Key Insights:

- Customers with **higher monthly charges** are more likely to churn, indicating that pricing strategy plays an important role in customer retention.
- Customers with **low tenure (new customers)** have a higher probability of leaving, showing that early engagement is critical.
- Long-term customers tend to remain loyal, suggesting that customer satisfaction improves over time.
- Certain features such as **Monthly Charges, Tenure, and Total Charges** have a strong influence on churn behavior.

- Machine learning models can effectively identify at-risk customers, enabling proactive retention strategies.

Business Impact: By identifying customers likely to churn, businesses can take targeted actions such as offering discounts, improving service quality, or providing personalized offers to retain customers and improve profitability.

15. BUSINESS RECOMMENDATIONS

Based on the analysis and model predictions, several actionable recommendations can be provided to improve customer retention and reduce churn.

Key Recommendations:

- Offer **discounts or loyalty benefits** to customers with high monthly charges to reduce churn risk.
- Focus on improving **onboarding experience for new customers**, as they are more likely to churn in the early stages.
- Implement **customer engagement programs** for short-tenure users to increase retention.
- Monitor customers with high predicted churn probability and provide **personalized retention offers**.
- Improve service quality and customer support to enhance satisfaction and long-term loyalty.

Business Strategy Impact: These recommendations help businesses proactively identify at-risk customers and take preventive actions, thereby improving customer retention rates and increasing overall revenue.

16. CONCLUSION

The Customer Churn Prediction project successfully demonstrates the application of machine learning techniques to analyze customer behavior and predict churn. The project involved data preprocessing, exploratory data analysis, feature engineering, model building, and evaluation to develop an effective predictive system.

Two machine learning models, Logistic Regression and Random Forest Classifier, were implemented and compared based on performance metrics such as accuracy, confusion matrix, and ROC-AUC score. Among these, the Random Forest model performed better in identifying customers who are likely to churn.

The analysis provided valuable insights into key factors influencing customer churn, such as monthly charges, tenure, and total charges. These insights can help businesses make data-driven decisions to improve customer retention strategies.

Overall, this project highlights the importance of predictive analytics in business decision-making and demonstrates how machine learning can be effectively used to solve real-world problems like customer

17. FUTURE SCOPE

The current Customer Churn Prediction system provides effective results using traditional machine learning models; however, there are several ways to enhance and extend this project in the future.

- Advanced machine learning and deep learning models such as XGBoost, Gradient Boosting, or Neural Networks can be implemented to improve prediction accuracy.
- Hyperparameter tuning techniques can be applied to optimize model performance.
- The system can be deployed as a real-time web application using frameworks like Streamlit or Flask for practical business use.
- Integration with live customer databases can enable real-time churn prediction.
- Additional customer behavioral and demographic data can be included to improve model robustness.

These improvements can make the system more scalable, accurate, and industry-ready, enabling businesses to take real-time proactive decisions for customer retention.

18. REFERENCES

1. Kaggle Dataset – Customer Churn Dataset
   https://www.kaggle.com/datasets (Customer Churn dataset source)
2. Pandas Documentation
   https://pandas.pydata.org/docs/
3. Scikit-learn Documentation
   https://scikit-learn.org/stable/documentation.html
4. Matplotlib Documentation
   https://matplotlib.org/stable/users/index.html
5. Microsoft Power BI Documentation
   https://learn.microsoft.com/en-us/power-bi/